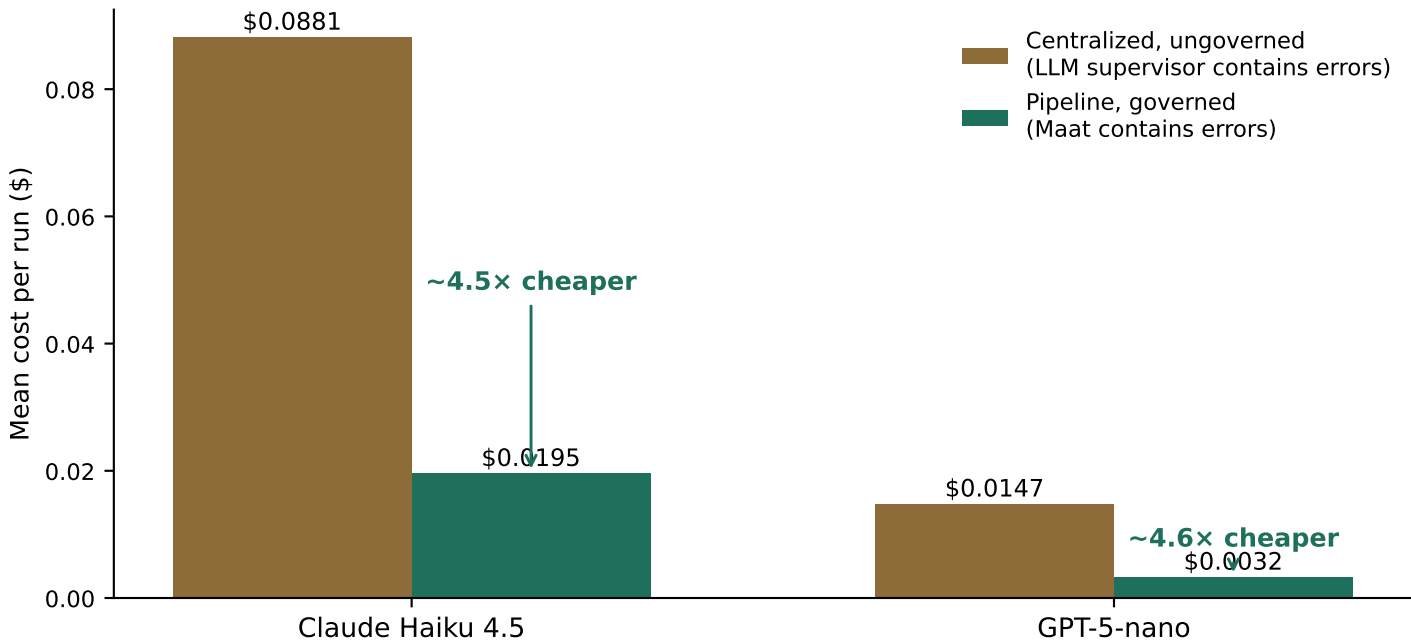


Figure T2-2. Same protection, without the supervisor



Both configurations achieved the identical outcome on both models: 0% defect propagation, €0 published harm.
The difference is how: an LLM orchestrator supervising every step (38.9% / 41.1% measured token overhead)
vs deterministic contract checks on the ordinary pipeline — at roughly one quarter of the cost.